JORDAN ELLIS

Machine Learning Engineer

Email: jordan.ellis@example.com | Phone: (555) 123-4567 | LinkedIn: linkedin.com/in/jordanellis-ml | GitHub: github.com/jordanellis-ml

PROFESSIONAL SUMMARY

Results-oriented Machine Learning Engineer with 5+ years of experience designing, building, and deploying scalable AI solutions into production environments. Adept at translating complex business challenges into robust statistical and deep learning models. Proven expertise in developing automated ML pipelines, fine-tuning large language models (LLMs), and optimizing model architectures for efficient deployment. Passionate about democratizing AI capabilities across cross-functional teams to drive product innovation.

CORE COMPETENCIES

- **Machine Learning & AI:** Deep Learning, NLP, Large Language Models (LLMs), Prompt Engineering, Data Augmentation.
- **MLOps & Infrastructure:** Automated ML Pipelines, Model Serving, CI/CD for ML, Model Compression, Quantization.
- **Languages & Frameworks:** Python, PyTorch, TensorFlow, Scikit-Learn, Hugging Face Transformers.
- **Tools & Platforms:** AWS (SageMaker, EC2, S3), Kubernetes, Docker, MLflow, Apache Kafka, Git.

PROFESSIONAL EXPERIENCE

Lead Machine Learning Engineer

January 2021 – Present

Nexus Data Solutions | San Francisco, CA

- **AI System Democratization:** Designed and built an internal AI platform that allowed cross-functional business teams (Marketing, Sales, and Operations) to leverage pre-trained predictive models, increasing cross-departmental AI adoption by 40%.
- **LLM Optimization & Fine-Tuning:** Developed fine-tuning frameworks for large language models specific to internal text-generation tasks. Implemented model compression and quantization techniques (e.g., INT8 quantization) that reduced the resource footprint of production LLMs by 35% with negligible loss in performance.
- **ML Infrastructure & Deployment:** Architected highly scalable ML infrastructure on AWS, integrating robust data storage, preprocessing, and model serving components to handle over 1M requests daily.
- **Automated Pipelines:** Built end-to-end automated pipelines using MLflow and Kubernetes for feature engineering, hyper-parameter tuning, and model evaluation, reducing the average experimentation-to-production lifecycle from 4 weeks to 8 days.
- **Governance & Collaboration:** Collaborated closely with product managers to define technical requirements from business challenges, establishing model governance protocols to ensure data privacy, fairness, and continuous performance monitoring.

Data Scientist / ML Developer
June 2018 – December 2020

Vanguard Analytics | Austin, TX

- **Statistical Solutions:** Partnered with the retail analytics team to design and implement deep learning forecasting models, improving inventory prediction accuracy by 22%.
- **Feature Engineering & Data:** Led the development of novel feature engineering and data augmentation workflows, significantly improving model robustness on imbalanced datasets.
- **Model Productionization:** Transitioned offline statistical models into robust REST APIs using FastAPI and Docker, ensuring seamless integration into the company's primary web applications.
- **Mentorship:** Guided junior analysts in transitioning from basic statistical analysis to foundational machine learning practices, promoting a culture of continuous learning and AI adoption.

EDUCATION

Master of Science in Computer Science (Focus on Artificial Intelligence)

University of Texas at Austin | Graduated May 2018

Bachelor of Science in Applied Mathematics

University of Michigan | Graduated May 2016